



COL333/671: Introduction to AI

Semester I, 2021

Adversarial Search

Rohan Paul

Outline



- Last Class
 - Local Search
- This Class
 - Adversarial Search
- Reference Material
 - AIMA Ch. 5 (Sec: 5.1-5.5)

Acknowledgement

These slides are intended for teaching purposes only. Some material has been used/adapted from web sources and from slides by Doina Precup, Dorsa Sadigh, Percy Liang, Mausam, Dan Klein, Anca Dragan, Nicholas Roy and others.

Game Playing and AI

- **Games: challenging decision-making problems**
 - Incorporate the state of the other agent in your decision-making. Leads to a vast number of possibilities.
 - Long duration of play. Win at the end.
 - Time limits: Do not have time to compute optimal solutions.



Games: Characteristics

- Axes:
 - Players: one, two or more.
 - Actions (moves): deterministic or stochastic
 - States: fully known or not.
- Zero-Sum Games
 - Adversarial: agents have opposite utilities (values on outcomes)

- **Core: contingency problem**
 - The opponent's move is **not** known ahead of time. A player must respond with a move for **every possible** opponent reply.
- **Output**
 - Calculate a **strategy (policy)** which recommends a move from each state.

Playing Tic-Tac-Toe: *Essentially a search problem!*

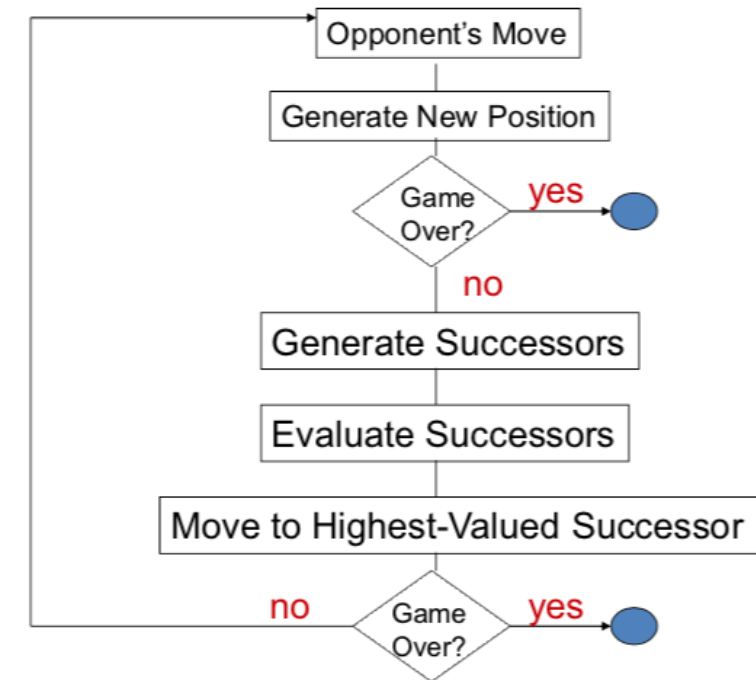
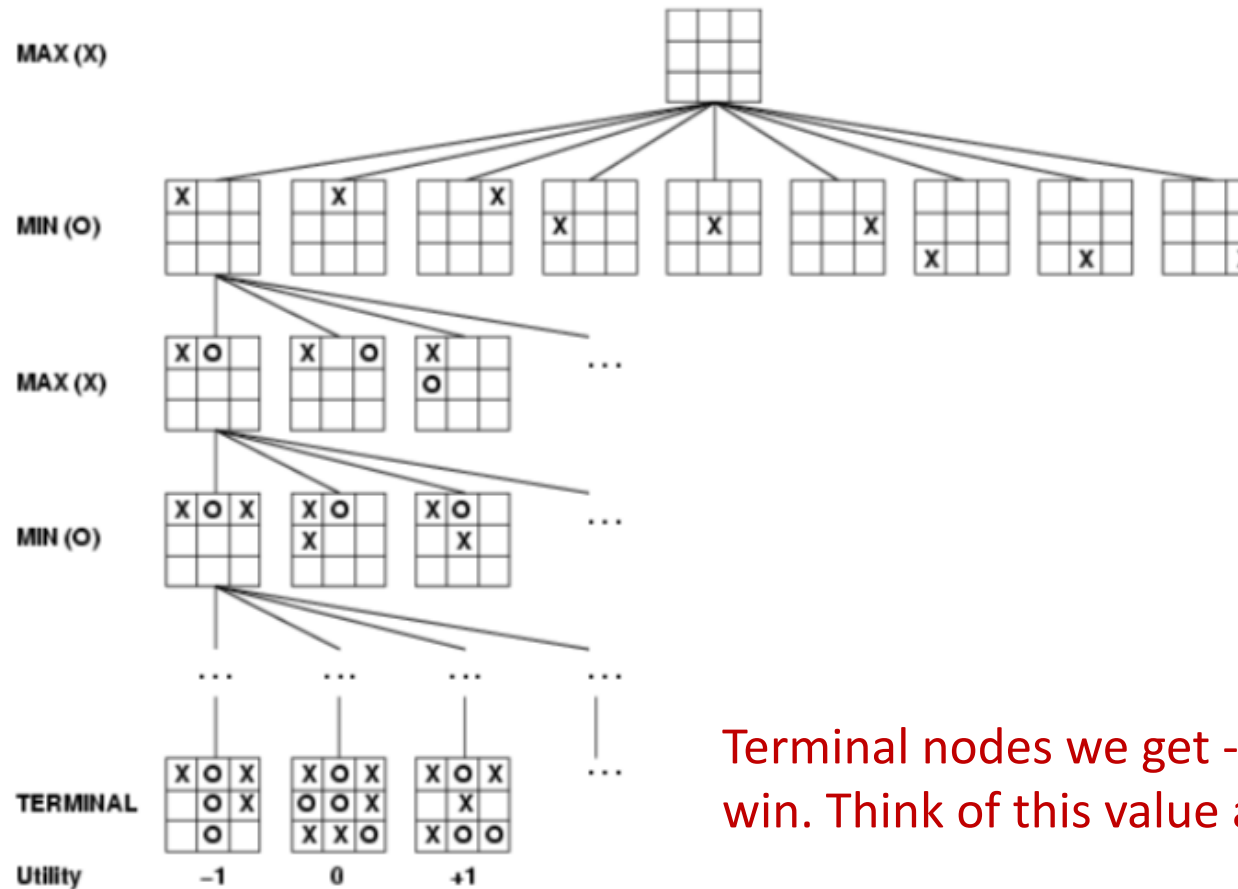
computer's
turn

opponent's
turn

computer's
turn


opponent's
turn

leaf nodes
are evaluated




Terminal nodes we get -1, 0 or 1 for loss, tie or win. Think of this value as a "utility" of a state.

Single-Agent Trees




Computing “utility” of states to decide actions

Value of a state:
The best achievable
outcome (utility)
from that state



Game Trees: Presence of an Adversary



The adversary's actions are not in our control. Plan as a contingency considering all possible actions taken by the adversary.


Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:


$$V(s) = \text{known}$$

Adversarial Search (Minimax)


- Consider a deterministic, zero-sum game
 - Tic-tac-toe, chess etc.
 - One player maximizes result and the other minimizes result.
- Minimax Search
 - Search the game tree for best moves.
 - Select optimal actions that move to a position with the highest minimax value.
 - What is the minimax value?
 - It is the best achievable utility against the optimal (rational) adversary.
 - Best achievable payoff against the best play by the adversary.

Minimax Algorithm

- Ply and Move
 - Move: when action taken by both players.
 - Ply: is a half move.
- Backed-up value
 - of a MAX-position: the value of the largest successor
 - of a MIN-position: the value of its smallest successor.
- Minimax algorithm
 - Search down the tree till the terminal nodes.
 - At the bottom level apply the utility function.
 - Back up the values up to the root along the search path (compute as per min and max nodes)
 - The root node selects the action.



Minimax Example



Minimax Implementation


def max-value(state):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{min-value}(\text{successor}))$

 return v



def min-value(state):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{max-value}(\text{successor}))$

 return v

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Implementation

```
def value(state):
```

if the state is a terminal state: return the state's utility

if the next agent is MAX: return `max-value(state)`

if the next agent is MIN: return `min-value(state)`

```
def max-value(state):
```

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

```
def min-value(state):
```

initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$

return v

Useful, when there are multiple adversaries.

Minimax Properties

- Completeness

- Yes

- Complexity

- Time: $O(b^m)$
 - Space: $O(bm)$
- Requires growing the tree till the terminal nodes.
- Not feasible in practice for a game like Chess.

- Chess:

- branching factor $b \approx 35$
 - game length $m \approx 100$
 - search space $b^m \approx 35^{100} \approx 10^{154}$

- The Universe:



- number of atoms $\approx 10^{78}$
 - age $\approx 10^{18}$ seconds
 - 10^8 moves/sec $\times 10^{78} \times 10^{18} = 10^{104}$

Minimax Properties

- Optimal


- If the adversary is playing optimally (i.e., giving us the min value)
 - Yes
- If the adversary is not playing optimally (i.e., not giving us the min value)
 - No. Why? It does not exploit the opponent's weakness against a suboptimal opponent).

You: Cricle. Opponent: Cross



If min returns 9? Or 100?


Necessary to examine all values in the tree?



Alpha-Beta Pruning: General Idea

- **General Configuration (MIN version)**


- Consider computing the MIN-VALUE at some node n , examining n 's children
- n 's estimate of the childrens' min is reducing.
- Who can use n 's value to make a choice? MAX
- Let a be the best value that MAX can get at any choice point along the current path from the root
- If the value at n becomes worse than a , MAX will not pick this option, so we can stop considering n 's other children (any further exploration of children will only reduce the value further)




Alpha-Beta Pruning: General Idea

- **General Configuration (MAX version)**


- Consider computing the MAX-VALUE at some node n , examining n 's children
- n 's estimate of the childrens' min is increasing.
- Who can use n 's value to make a choice? MIN
- Let b be the lowest (best) value that MIN can get at any choice point along the current path from the root
- If the value at n becomes higher than b , MIN will not pick this option, so we can stop considering n 's other children (any further exploration of children will only increase the value further)



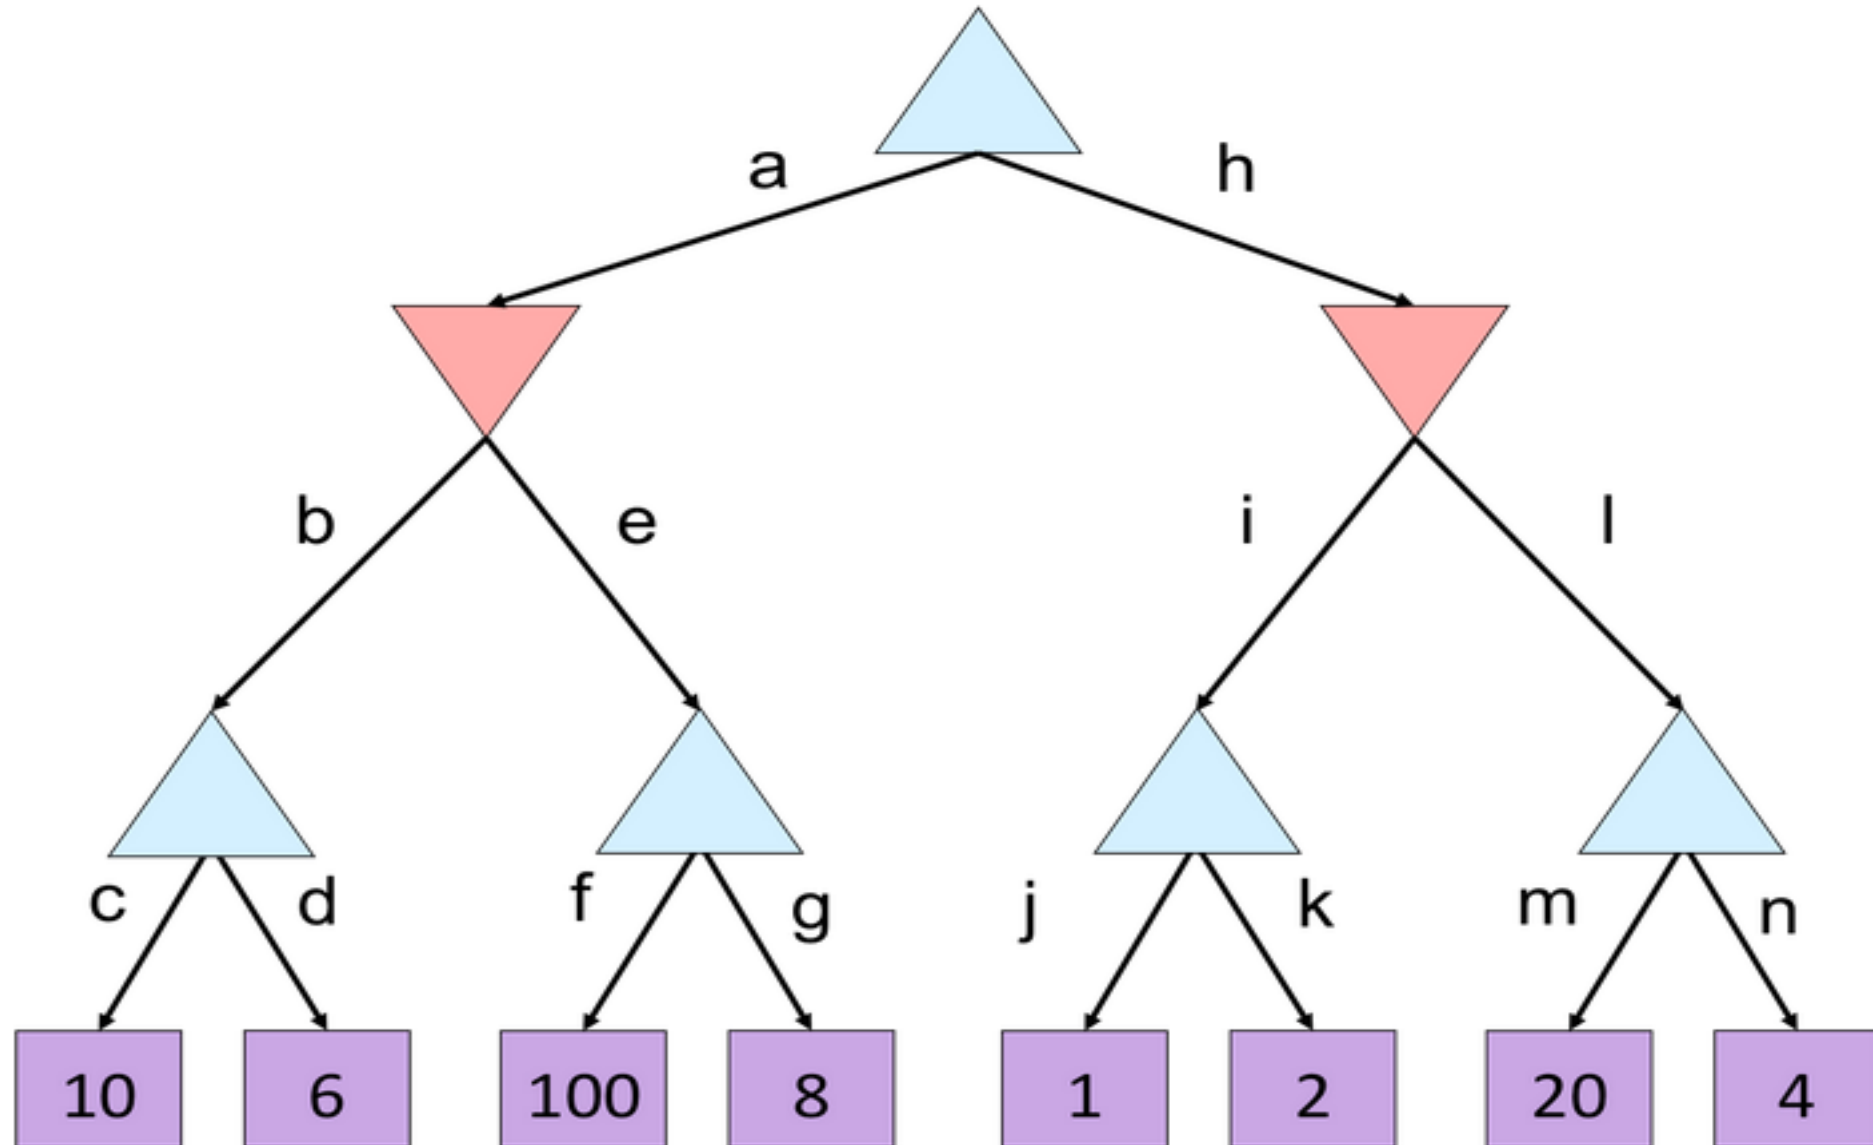
Pruning: Example




Pruning: Example



Pruning: Example



Pruning: Example



Alpha-Beta Implementation

α : MAX's best option on path to root
 β : MIN's best option on path to root

def max-value(state, α , β):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

if $v \geq \beta$ return v

$\alpha = \max(\alpha, v)$

return v

def min-value(state, α , β):

initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

if $v \leq \alpha$ return v


$\beta = \min(\beta, v)$

return v


Alpha-Beta Pruning - Properties

1. Pruning has **no effect** on the minimax value at the root.
 - Pruning does not affect the final action selected at the root.
2. A form of **meta-reasoning** (computing what to compute)
 - Eliminates nodes that are irrelevant for the final decision.

Alpha-Beta Pruning – Order of nodes matters



Alpha-Beta Pruning – Order of nodes matters




Alpha-Beta Pruning - Properties


1. Pruning has **no effect** on the minimax value at the root.
 - Pruning does not affect the final action selected at the root.
2. A form of **meta-reasoning** (computing what to compute)
 - Eliminates nodes that are irrelevant for the final decision.
3. The alpha-beta search cuts the largest amount off the tree when we examine the **best move first**
 - However, best moves are typically **not** known. Need to make estimates.

Alpha-Beta Pruning – Order of nodes matters

- Bad: Worst moves encountered first



- Good: Good moves ordered first



Ordering moves with good moves first can benefit alpha-beta pruning.

Alpha-Beta Pruning – $O(b^{m/2})$

Let $T(m)$ be time complexity of search for depth m

Normally:

$$T(m) = b.T(m-1) + c \rightarrow T(m) = O(b^m)$$

With ideal α - β pruning:

$$T(m) = T(m-1) + (b-1)T(m-2) + c \rightarrow T(m) = O(b^{m/2})$$

We are cutting off the branching at every other level.

Alpha-Beta Pruning - Properties

1. Pruning has **no effect** on the minimax value at the root.
 - Pruning does not affect the final action selected at the root.
2. A form of **meta-reasoning** (computing what to compute)
 - Eliminates nodes that are irrelevant for the final decision.
3. The alpha-beta search cuts the largest amount off the tree when we examine the **best move first**
 - Problem: However, best moves are typically **not** known.
 - Solution: Perform iterative deepening search and evaluate the states.
4. Time Complexity
 - **Best ordering** - $O(b^{m/2})$. Can double the search depth for the same resources.
 - On average – $O(b^{3m/4})$ if we expect to find the min or max after $b/2$ expansions.

Minimax for Chess

- Chess:
 - branching factor $b \approx 35$
 - game length $m \approx 100$
 - search space $b^m \approx 35^{100} \approx 10^{154}$
- The Universe:
 - number of atoms $\approx 10^{78}$
 - age $\approx 10^{18}$ seconds
 - 10^8 moves/sec $\times 10^{78} \times 10^{18} = 10^{104}$

Alpha-Beta for Chess

- Chess:
 - branching factor $b \approx 35$
 - game length $m \approx 100$
 - search space $b^{m/2} \approx 35^{50} \approx 10^{77}$

Cutting-off Search

- Problem:


- Minimax search: full tree till the terminal nodes.
- Alpha-beta prunes the tree but still searches till the terminal nodes.
- Still difficult to search till the leaves.

- Solution:

- Depth-limited Search (H-Minimax)
- Search only to a limited depth (cutoff) in the tree
- Replace the terminal utilities with an evaluation function for non-terminal positions


H-MINIMAX(s, d) =

$$\begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MIN.} \end{cases}$$



Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search.
- Estimate the chances of winning.




- Ideal function: returns the actual **minimax** value of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- e.g. $f_i(s)$ = (number of pieces of type i), each weight w_i etc.

Incorporating Chance: Expectimax Search

- Till now
 - Assumed that the opponent provides us with the *worst-case* outcome.
- Incorporate a notion of chance
 - Include chance nodes
 - Unpredictable opponents: the ghosts move randomly in Pacman
 - Explicit randomness: rolling dice
- Computing values at nodes
 - Not worst-case (minimax) outcomes
 - Reflect average-case (expectimax) outcomes
- **Expectimax search:**
 - Compute the average score under optimal play
 - Max nodes as in minimax search
 - At chance nodes the outcome is uncertain
 - Calculate **expected utilities**: weighted average (expectation) of children



Expectimax Search

```
def value(state):
```

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is EXP: return exp-value(state)

```
def max-value(state):
```

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

```
def exp-value(state):
```


initialize $v = 0$

for each successor of state:

$p = \text{probability}(\text{successor})$


$v += p * \text{value}(\text{successor})$

return v





Expectimax Search

```
def exp-value(state):  
    initialize v = 0  
    for each successor of state:  
        p = probability(successor)  
        v += p * value(successor)  
    return v
```

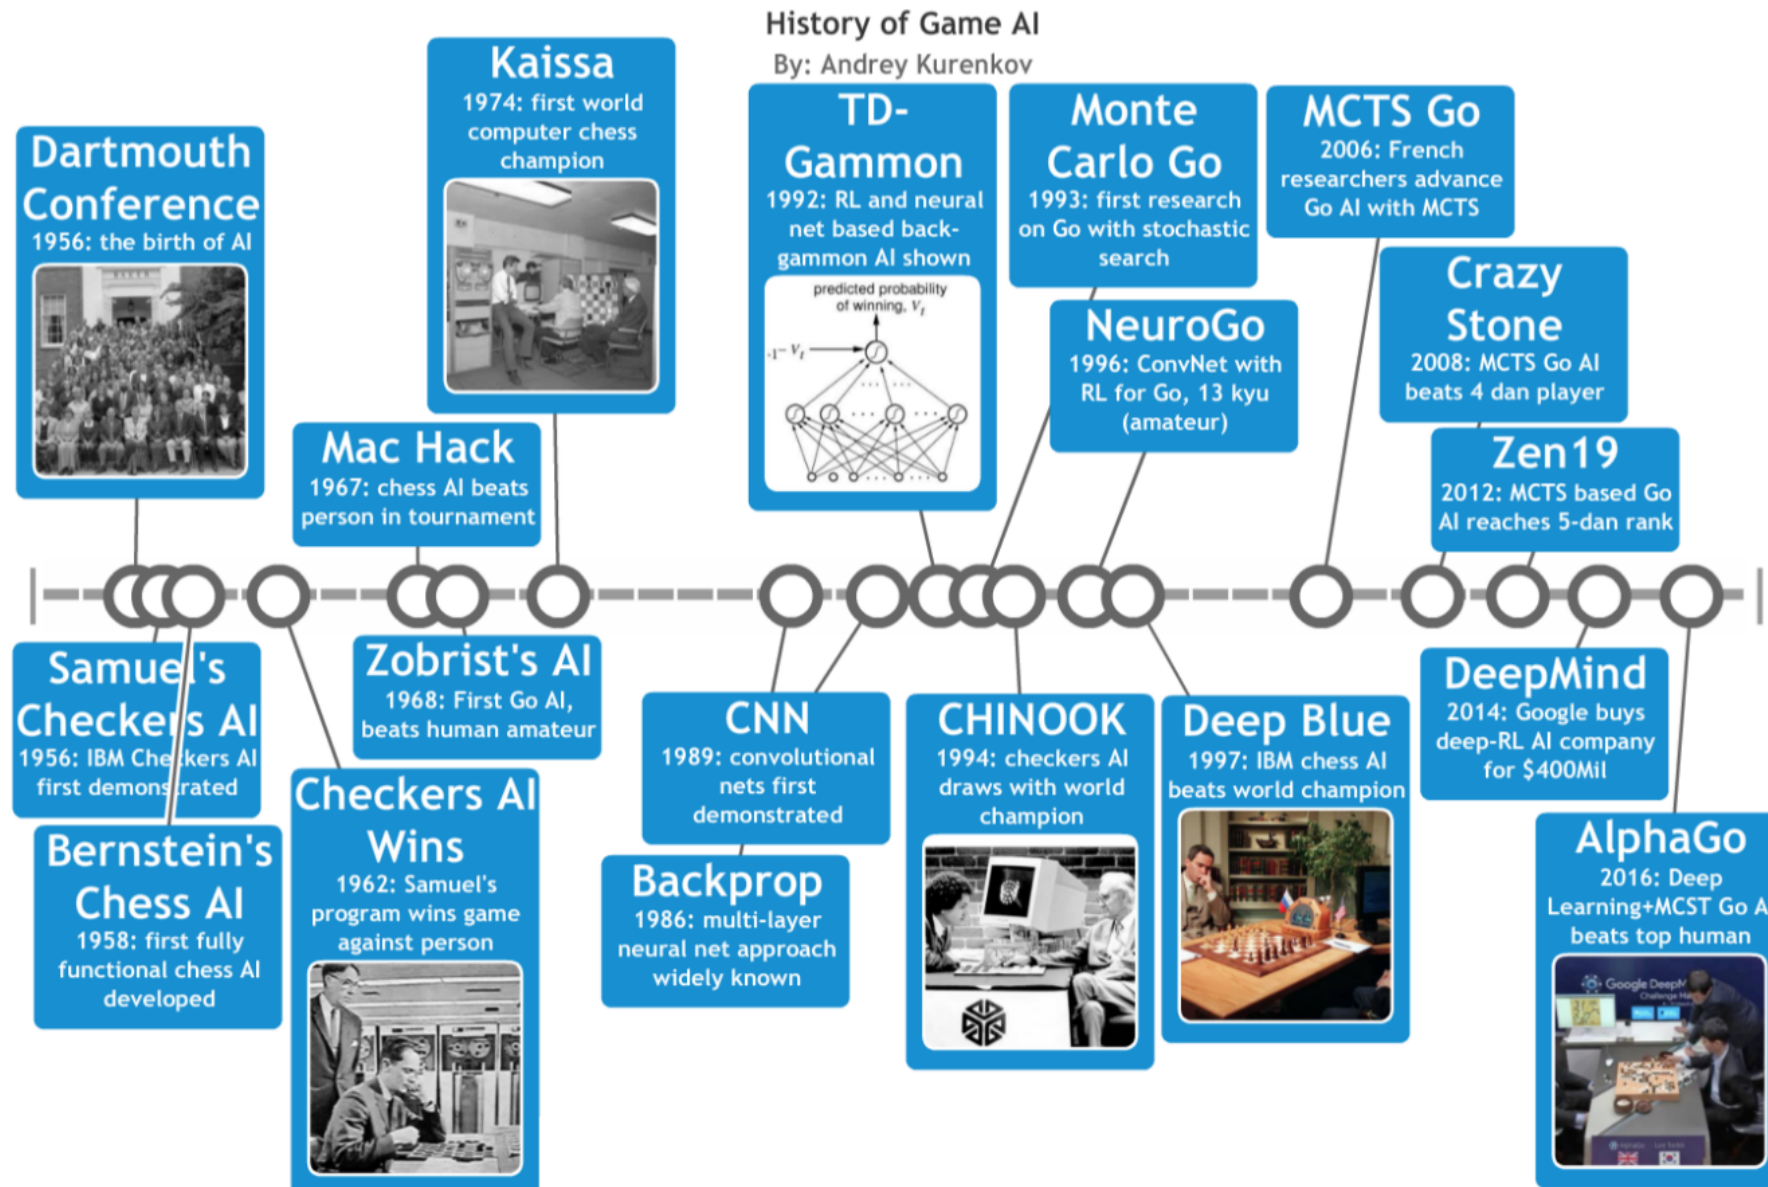


$$v = (1/2) (8) + (1/3) (24) + (1/6) (-12) = 10$$

Expectimax Search



Can we perform pruning?



“Games are to AI as grand prix is to automobile design”
Games viewed as an indicator of intelligence.